

PyTorch Conference 2022 Tech Report

Won Jong Jeon
Software Lab

Overview

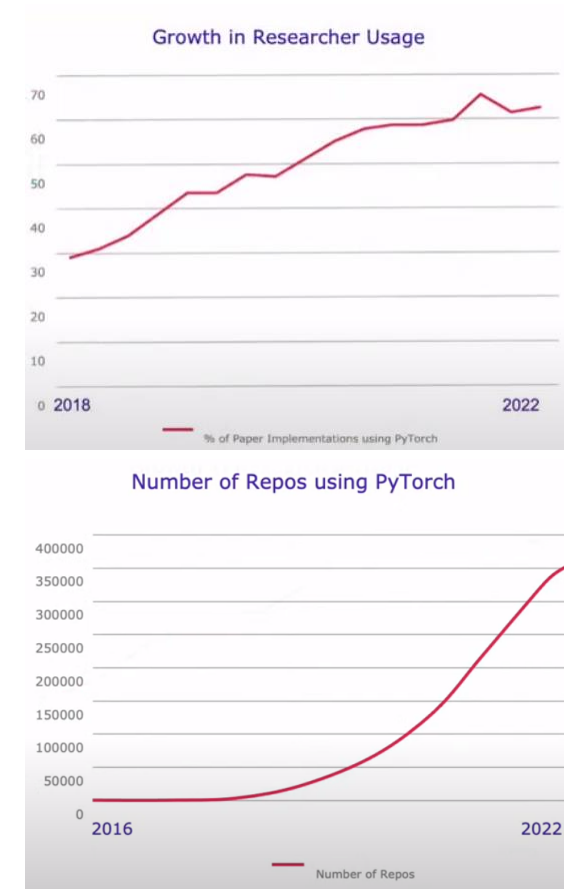
- When: December 2, 2022
- Where: New Orleans, LA, USA (virtual option available)
- Co-located with NeuralPS 2022
- Covering new software releases on PyTorch, use cases in academia and industry, as well as ML/DL development and production trends.
- Video stream is available at <https://youtu.be/vbtGZL7IrAw>.

Schedule

- 8-9am: Registration/check-in
- 9-11:20am: Keynote & technical talks (by Meta AI)
- 11:30am-1pm: Lunch
- 1-3pm: Poster session & breakouts
- 3-4pm: Community/partner talks
- 4-5pm: Panel discussion

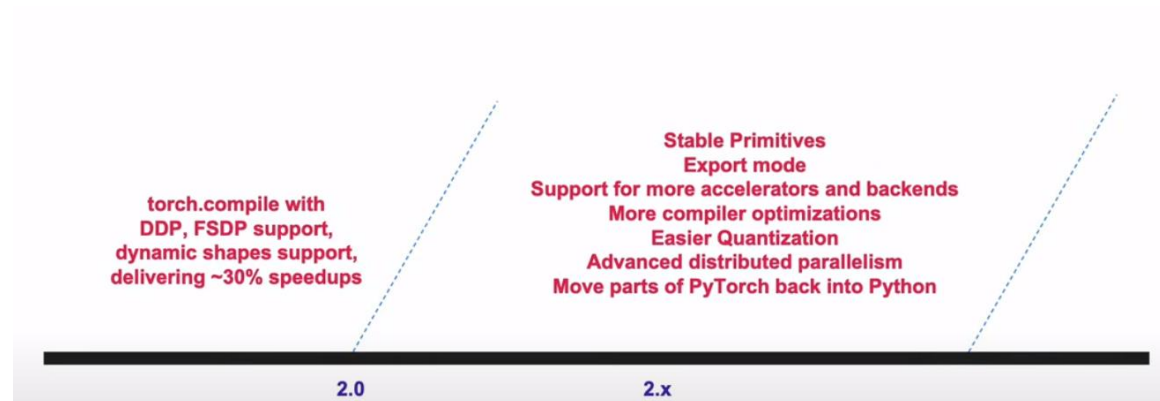
Introduction

- Rapid growth in research usage
- Rapid increase in # of repos using PyTorch
- Industrial usage
 - >44k LinkedIn professionals, >2,500 jobs, 50% increase in PyTorch professionals this year
- Top organizations contributing to PyTorch
 - Meta, Microsoft, Nvidia, Intel, Google, Quansight, AMD, AWS, IBM
- PyTorch Foundation
 - Under Linux Foundation (LF)
 - For technical autonomy, business governance



Plan/roadmap of PyTorch 2.0 and beyond

- For ML scientists
 - 30%+ training speedups, lower memory usage with no changes to code or workflow
- For compiler/hardware engineer
 - Dramatically easier to write a PyTorch backend
- For large-scale DL projects
 - State of the art distributed capabilities
- For code contributor
 - Substantially more PyTorch is written in Python



Plan/roadmap

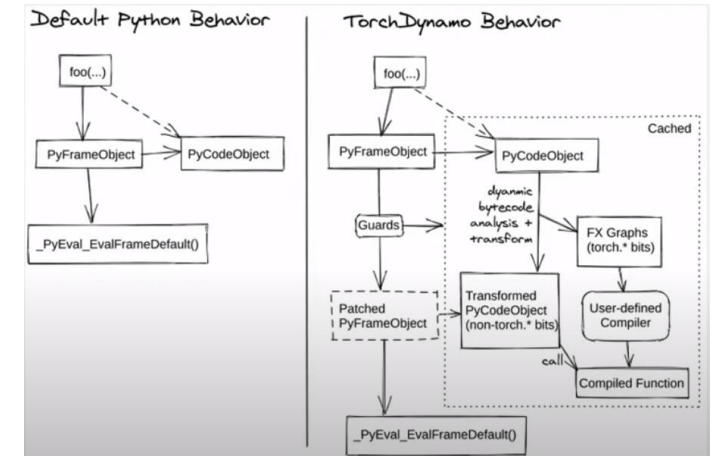
- 2.0 is fully backward compatible.
- Sets of benchmarks FP16 and FP32
 - timm (Python image models) <https://github.com/rwightman/pytorch-image-models>
 - TorchBench <https://github.com/pytorch/benchmark>
 - HuggingFace transformer benchmark <https://huggingface.co/docs/transformers/benchmarks>
- Dynamic shape and distributed support
- Robust to correctness and accuracy
 - AMP, FP16 / FP32
- Make PyTorch faster
 - Kernel fusion
 - Out-of-order execution
 - Automatic work placement (in multi-node multi-GPU environment)

Plan/road map

- TorchDynamo
 - To capture dynamic graphs without compromising user-experience?
 - TorchDynamo rewrite into blocks of graphs.
- AOT Autograd
 - https://pytorch.org/functorch/stable/notebooks/aot_autograd_optimizations.html
- Simplifying operator space
 - From 2000+ ops to ~250 primitive operators
- TorchInductor
 - Graph compilation, powered by OpenAI Triton (Python DSL for writing parallel code) <https://triton-lang.org>
 - Compiler written in Python
 - Support for CPU and GPU (Volta and Ampere)
 - Support for your own backend, nvFuser, TVM, XLA, AITemplate, TensorRT

TorchDynamo

- Graph capture fundamentally shift the efficiency of PyTorch.
- Features
 - Partial graph capture
 - Ability to skip unwanted parts of eager
 - Guarded graphs
 - Ability to check if captured graph is valid for execution
 - Just-in-time recapture
 - Recapture a graph if captured graph is invalid for execution
- TorchDynamo vs. TorchScript frontend
 - TorchDynamo does not require changing models
 - TorchDynamo reliably captures backward graphs (works for training)
- Status
 - Tested with 7k+ Github models, 20+ inference backends, 1+ training backends
 - 30+% geomean speedup (FP32, AMP, on A100 GPU)
- Current recommendations on graph capture
 - For training: all (but XLA or TPU) → Dynamo, Export to XLA or TPU → Lazy Tensor
 - For inference: embedded → TorchScript, non-embedded → TorchScript or torch.fx
 - For human-in-the-loop tool → torch.fx



TorchInductor

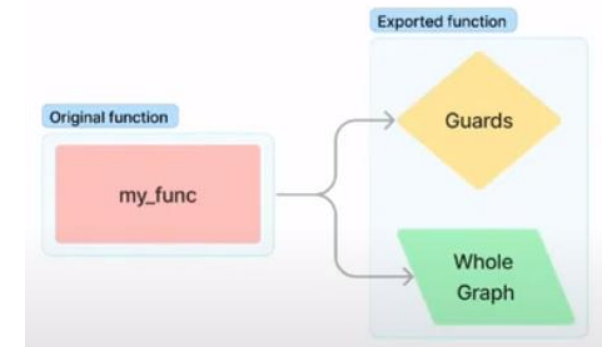
- A new compiler backend for PyTorch
- Principles
 - Similar abstractions to PyTorch eager
 - Written in Python, generates Triton and C++
 - Early focus on supporting a variety of operators, hardware, and optimization
- Technologies
 - Define-by-run loop-level IR: direct use of Python functions in IR definitions
 - Dynamic shape & strides: using SymPy symbolic math library (<https://www.sympy.org>) to reason about shapes, indexing, and managing guards.
 - Reuse state of the art languages: Triton for GPUs, C++/OpenMP for CPUs
- https://github.com/pytorch/pytorch/tree/master/torch/_dynamo

AOT Autograd / PrimTorch Decomposes into smaller operator set Capture forwards + backwards Some inductor specific decomp included in this step	Inductor Graph Lowerings Remove views, broadcasting, and simplify indexing Rematerialize vs reuse decisions Layout tuning and optimization Loop order	Inductor Scheduling Horizontal / vertical fusion decisions Reduction fusions Tiling Memory planning and buffer reuse In-place memory buffers Autotuning / kernel selection	Wrapper Codegen Outer code that calls kernels and allocates memory (Replaces interpreter) Backend Codegen Triton C++
--	--	---	---

```
def inner_fn(index: List[sympy.Expr]):  
    i1, i0 = index  
    tmp0 = ops.load("x", i1 + i0*size1)  
    tmp1 = ops.load("x", 2*size1 + i0)  
    return ops.add(tmp0, tmp1)  
  
torchinductor.ir.Pointwise(  
    device=torch.device("cuda"),  
    dtype=torch.float32,  
    inner_fn=inner_fn,  
    ranges=[size0, size1],  
)
```

PyTorch 2.0 Export


- Whole-graph export API
- Standardized IR and operator set
 - Primitive operator set (from 2000+ ops to ~250)
 - Consolidated compiler infrastructure: provide common infra for users to process graphs , vendor common passes in reusable form
- TorchDynamo export mode
 - `torch._dynamo.export(my_func, input)`
- Status
 - Significant functionality already in repo tree
 - Will be stabilized in 2.x series release




Dynamic shape support in PyTorch 2.0

- Challenge
 - Python code, PyTorch C++, and kernels in eager mode can depend on shapes.
 - Cache in PyTorch 2.0 needs to know about shapes.
- Solution
 - Transition from old concrete static shapes to symbolic shapes
 - Leverages SymPy to build rich information about the shapes in the program
 - Allows TorchInductor to generate efficient code for different shapes without recompilation
 - Deep integration of symbolic shapes into PyTorch code components
- Performance enhancement in execution time and compilation time
- Additional benefits
 - Shape-checking of functions
 - Analyzing FLOPs of neural network symbolically

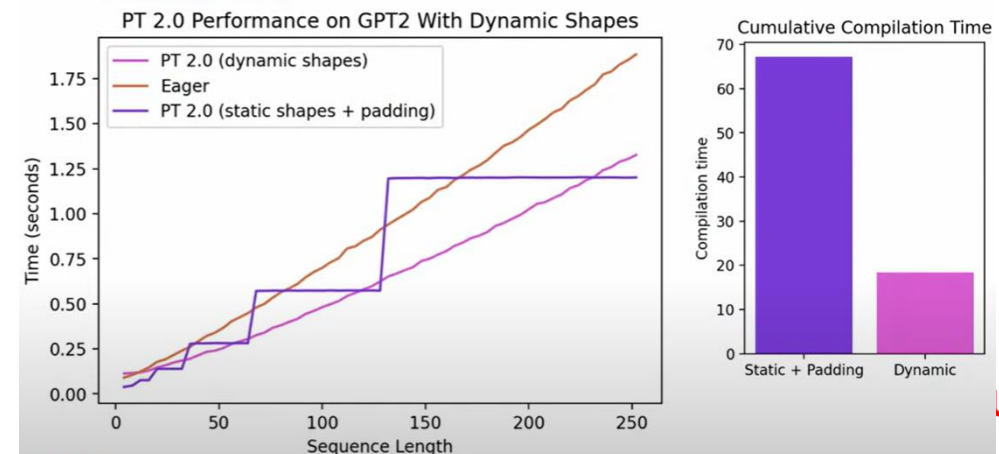
Eager



PyTorch 2.0

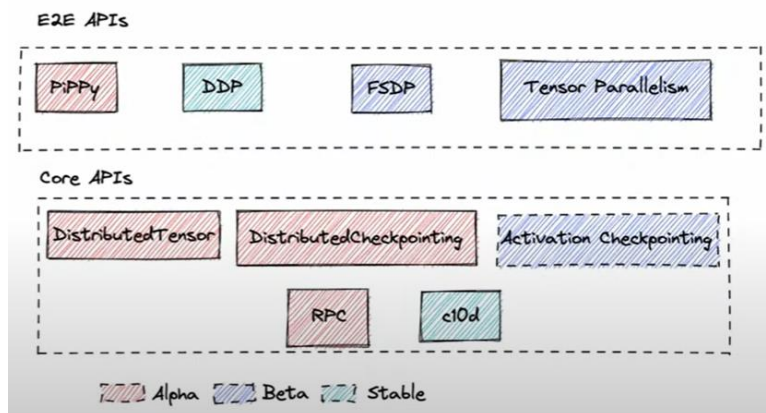


Performance Results



PyTorch Distributed


- DDP (Distributed Data Parallel) and FSDP (Fully Shared Data Parallel) for data parallel training
 - DDP working by replicating parameters
 - FSDP sharding parameters and using allgather to ensure replicated parameters on each device
 - Recent improvement including activation checkpointing, mixed precision, tackling OOM issues
 - Dynamo optimizes DDP wrapped modules
- PiPPy
 - A cross host pipeline parallelism API, enabling automatic splitting of model using torch.fx and 2D parallelism with DDP
 - Recent improvement in deferred initialization, automated model split APIs (for user-guided splits)
- Tensor Parallelism Modules (Dtensor)
 - Dtensor is a subclass of tensor, providing a flexible annotation for both Shared and Replicated Tensors used in Tensor Parallelism.
 - Tensor Parallelism can be used together with FSDP in 2D parallel.



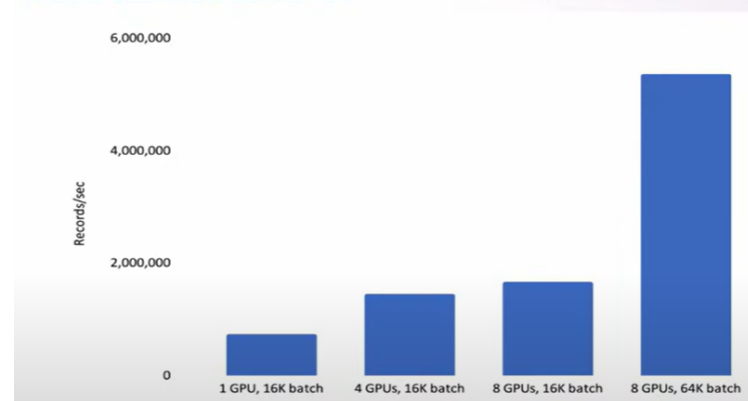
TorchRec

- PyTorch domain library for large-scale recommender systems
 - Domain specific: TorchRec's custom PyTorch modules are built for recommender systems and optimized for distributed run-time environments.
 - Scalable: TorchRec leverages Model Parallelism and automatically adjusts as authors scale from 1 to N devices.
 - Performant: TorchRec's performance optimizations are born from research and built for production.
- Module-based model parallelism
 - TorchRec prepares PyTorch models for distributed training or inference.
- Additional features
 - Batched embeddings, fused optimizers, jagged tensors, hierarchical sharding, input batch pipelining, collectives quantization, embedded quantization, automated planning, HBM/DDR caching, ...
- Performance enhancements
 - HBM (High Bandwidth Memory) and UVM (Unified Virtual Memory) optimizations
 - Sharded DLRM model over multiple GPUs

DLRM MODEL SINGLE RANK BENCHMARKS



SHARDED DLRM MODEL BENCHMARKS



torch::deploy (MultiPy)

- Running multiple Python interpreters inside a **single process**
- Running eager mode PyTorch models in production
- No GIL (Global Interpreter Lock)
- C++ library: wrapping Python object into C++
- Linux x86_64 (arm64 support as well)
- Currently in beta
- No modifications is needed, no tracing/scripting is needed
- Share the same model across multiple Python interpreters.
- Shared backend: libtorch/aten backend is the same, no extra copies of models
- No process boundaries: simplify production stack by eliminating data transfer between processes

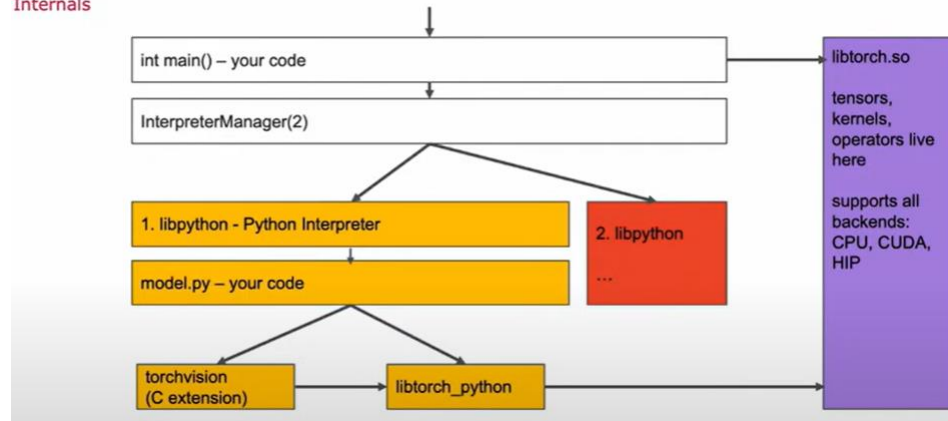
```
#include <torch/extension.h>
...
// 1. Create two interpreters
multipy::InterpreterManager m(2);

// 2. Acquire an interpreter
auto i = m.acquireOne();

// 3. Construct your model in Python
auto model = i.global("torch.nn.Conv2d")({{2, 2, 1}});

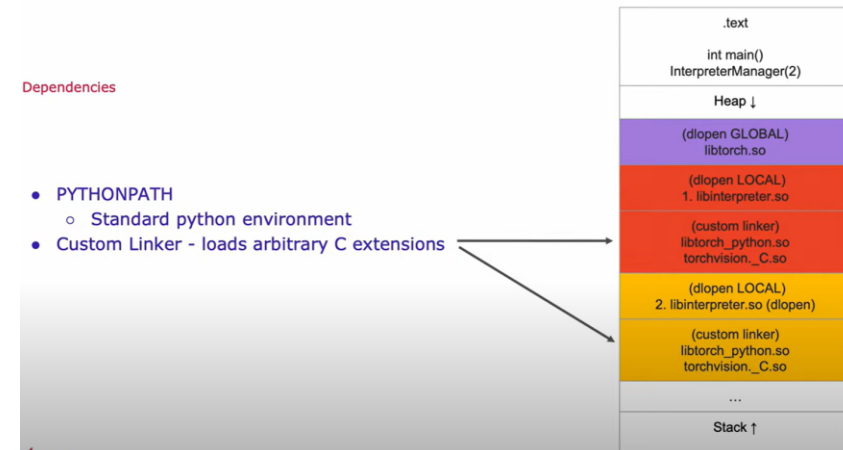
// 4. Run it
std::Tensor output = model->forward({{1, 1, 1, 1}});
return output.toTensor();
```

Internals



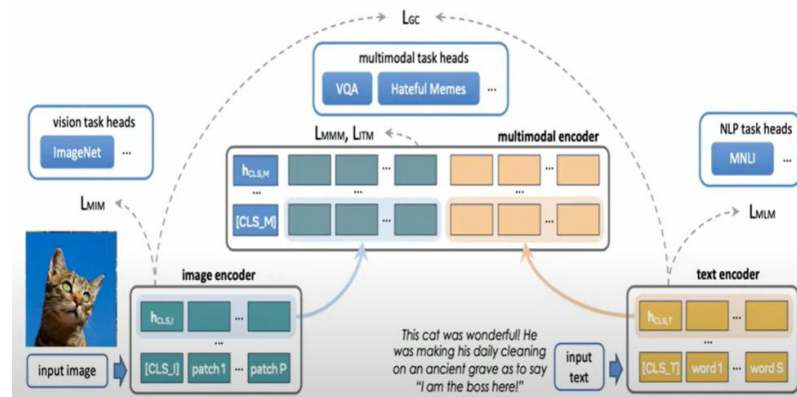
Dependencies

- PYTHONPATH
 - Standard python environment
- Custom Linker - loads arbitrary C extensions



TorchMultimodal (beta)

- Tasks that can understand different types of input (e.g., textual, visual , audio)
 - Visual question answering, text to image / video retrieval, text to image generation
- Can use the understanding to generate outputs
 - Content understanding, integrity classifiers, self-driving cars
- Core principles
 - Modularity: modular low-level components can be used independently
 - Interoperability: enable to plug in components from other PyTorch libraries with minimal effort
 - Extensibility: task agnostic models can be extended with any task-specific layers
- Example offerings
 - Building blocks: vector quantized VAE, contrastive loss with temperature
 - End-to-end models: CLIP (Contrastive Language-Image Pretraining), MDETR (Modulated DETR)
 - Example scripts: Omnivore (single model for many visual modalities), FLAVA (Foundational Language And Vision Alignment Model)



FLAVA demo

TorchRL

- Reinforcement learning and control library for PyTorch
 - Modular enough to easily swap between components
 - Syntax is familiar to RL practitioners
 - Aiming at PyTorch, but optionally support other libraries (e.g., gym/gymnasium, dm_control, habitat-lib, Jumanji, SMAC, dm_lab*, unity*, brax*)
- Efficiency
 - Efficient distributed replay buffer
 - Vectorized environments and transforms, fully operational on device
 - Vectorized advantage computation (10-100x faster)
- Modularity
 - Generic module class with few levels of abstraction (easy to hack)
 - Environments, modules, models, losses, etc. are supposed to be re-used across frameworks
 - MBRL/MFRL, off-/on-policy
 - Multi/single agent, offline RL, multi-task, distributed, meta-RL, MCTS, MBRL
- TensorDict
 - A new tensor container that allows to abstract away the irrelevant parts of each module.

On-device ML in PyTorch

- Challenges
 - More software platforms (OS, hardware)
 - More devices and hardware heterogeneity
 - Mixed compute unit execution (e.g., DSP+NPU)
 - More metrics to optimize for (e.g., power)
 - Increased model variety and complexity
 - Number and variety of experts involved
- Principles
 - Maintain PyTorch authoring semantics: clear programming model
 - Provide an extensible software stack: easy customization of the program
 - Supply out-of-the-box foundational components: portable and efficient runtime, compiler-based backends
 - Offer high developer productivity: developer SDK, documentations, examples, ...
- Roadmap
 - 2019 PT1.3 experimental for Android/iOS
 - 2020 PT1.4 performance improvements
 - 2021 PT1.9 75% binary size reduction
 - Beta release coming late summer next year